Anmol Saluja

Data Analytics · Business Analysis · Business Intelligence

Toronto, ON
anmolsaluja045@gmail.com
LinkedIn · GitHub

Business analytics professional with UBC Sauder training and experience in fintech delivery and growth strategy. Uses SQL, Python, R, Tableau, Power BI and Excel for data analysis, reporting and stakeholder communication. Built and deployed Laptop Intelligence: an AI research agent combining ETL, semantic search, competitor benchmarking, NLP sentiment, retrieval-augmented generation and data-backed visualizations.

TECHNICAL SKILLS

Analytics & programming: SQL · Python · R · Advanced Excel · pandas · ETL / Parquet

Business intelligence: Tableau · Power BI · KPI reporting · Data visualization · Stakeholder communication

Applied AI & modelling: Prompt engineering · Agentic AI · LLMs · RAG · Precision / recall / F1 · Confusion matrices · RAG answer audits · XGBoost · Regression

PROFESSIONAL & EXPERIENTIAL LEARNING

UBC Experiential Learning

May – Aug 2026

Business Analytics Internship · Vancouver, BC

- Built Laptop Intelligence, a live AI research agent for competitor benchmarking, pricing analytics and customer insights across 25,841 laptop listings, 456,395 reviews and 294 brands.
- Streamed 27.9 GB of Amazon data into pandas and Parquet datasets. Combined MiniLM embeddings with hardware features; the project guard audit reported 100% segment consistency across 5,000 recommendations.
- Evaluated DistilBERT sentiment against a 6,000-review star-label benchmark at 89.0% accuracy, identifying customer support and battery feedback as improvement themes.
- Integrated a Groq-hosted GPT-OSS-120B RAG assistant with intent routing, evidence citations, answer checks and natural-language chart tools. Deployed the FastAPI / JavaScript application using Docker, Caddy and HTTPS on Oracle Cloud.

FIS Global

Apr 2022 – Apr 2023

Engineer · India

- Coordinated cross-functional teams implementing large financial technology projects and supported collaboration across business and delivery teams.
- Designed and implemented an interactive Power BI dashboard that reduced report generation time from 25 minutes to 3 minutes.
- Used reporting and visualization to improve transparency and support stakeholder decision-making.

Conquering My Own Everest / GoGhumantu

Jan 2021 – Mar 2022

Growth Manager · India

- Developed and executed a growth strategy that increased online engagement by 100%.
- Led a digital outreach campaign that increased customer acquisition from 2% to 10%.
- Developed subscription pricing recommendations using demand, market trends and revenue drivers.

EDUCATION

UBC Sauder School of Business

2026

Master of Business Analytics · Vancouver, BC

Degree conferral: November 2026

Sharda University

2020

Master of Business Administration · Noida, India

SELECTED ACADEMIC PROJECTS

Yelp restaurant & customer analytics

Academic project

7,493 restaurants · Tableau · KPI dashboards · Geospatial analysis · Cohort analysis

Built a Tableau restaurant overview and review / user storyboards. Combined cuisine filters, map-based selection, popularity rankings, rating distributions and time-series views to explore Greater Vancouver's restaurant landscape. The analysis separates a sharp review-volume decline around March 2020 from comparatively stable average ratings, and highlights concentrated reviewer influence. Review counts measure platform activity, not sales or customer footfall.

Pricing strategy & demand elasticity

Academic project

86 stores · 3 products · R · Log–log regression · Price elasticity · Marketing analytics

Estimated own- and cross-price elasticities for Tide 128 oz, Tide 64 oz and Wisk 64 oz using weekly store-level grocery data. Excluded promotional weeks and controlled for store effects and a weekly time trend. All three reported own-price elasticities are below −1. Wisk 64 oz is most price-sensitive (≈ −2.98); Tide 64 oz is least sensitive (≈ −1.49). These are observational estimates, not measured effects of a pricing experiment.

PetFinder adoption-speed prediction

Feb 2026

14,993 pet listings · Python · scikit-learn · XGBoost · Text features · Model evaluation

Compared five classifiers for adoption-speed prediction using pet characteristics and listing information. Merged the two fastest original categories into four classes and used stratified train / validation / test splits. The final report records 42.85% test accuracy and 0.4175 macro-F1 for XGBoost + TF-IDF on 2,999 held-out listings. Intermediate speed classes remain difficult; listing-improvement recommendations are hypotheses, not measured adoption gains.

Featured project: Laptop Intelligence — live application · Source code